

<https://www.pyimagesearch.com/2017/10/16/raspberry-pi-deep-learning-object-detection-with-opencv/>

# Raspberry Pi: Deep learning object detection with OpenCV

by [Adrian Rosebrock](#) on October 16, 2017 in [Deep Learning](#), [Object Detection](#), [OpenCV](#), [Raspberry Pi](#)

A few weeks ago I demonstrated how to perform [real-time object detection using deep learning and OpenCV](#) on a standard laptop/desktop.

After the post was published I received a number of emails from PyImageSearch readers who were curious if the Raspberry Pi could *also* be used for real-time object detection.

The short answer is “kind of” ...

*...but only if you set your expectations accordingly.*

Even when applying our [optimized OpenCV + Raspberry Pi install](#) the Pi is only capable of **getting up to ~0.9 frames per second** when applying deep learning for object detection with Python and OpenCV.

Is that fast enough?

Well, that depends on your application.

If you’ re attempting to detect objects that are quickly moving through your field of view, likely not.

But if you’ re monitoring a low traffic environment with slower moving objects, the Raspberry Pi could indeed be fast enough.

In the remainder of today’ s blog post we’ ll be reviewing two methods to perform deep learning-based object detection on the Raspberry Pi.

**Looking for the source code to this post?**

[Jump right to the downloads section.](#)

# Raspberry Pi: Deep learning object detection with OpenCV

Today's blog post is broken down into two parts.

In the first part, we'll benchmark the Raspberry Pi for real-time object detection using OpenCV and Python. This benchmark will come from the exact code we used for our [laptop/desktop deep learning object detector](#) from a few weeks ago.

I'll then demonstrate how to use multiprocessing to create an alternate method to object detection using the Raspberry Pi. This method may or may not be useful for your particular application, but at the very least it will give you an idea on different methods to approach the problem.

## Object detection and OpenCV benchmark on the Raspberry Pi

The code we'll discuss in this section is identical to our previous post on [Real-time object detection with deep learning and OpenCV](#); therefore, I will not be reviewing the code exhaustively.

For a deep dive into the code, please see the [original post](#).

Instead, we'll simply be using this code to benchmark the Raspberry Pi for deep learning-based object detection.

To get started, open up a new file, name it `real_time_object_detection.py`, and insert the following code:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
1 # import the necessary packages
2 from imutils.video import VideoStream
3 from imutils.video import FPS
4 import numpy as np
5 import argparse
6 import imutils
7 import time
8 import cv2
```

We then need to parse our command line arguments:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
10 # construct the argument parse and parse the arguments
11 ap = argparse.ArgumentParser()
12 ap.add_argument("-p", "--prototxt", required=True,
```

```

13         help="path to Caffe 'deploy' prototxt file")
14     ap.add_argument("-m", "--model", required=True,
15         help="path to Caffe pre-trained model")
16     ap.add_argument("-c", "--confidence", type=float, default=0.2,
17         help="minimum probability to filter weak detections")
18     args = vars(ap.parse_args())

```

Followed by performing some initializations:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```

20     # initialize the list of class labels MobileNet SSD was trained to
21     # detect, then generate a set of bounding box colors for each class
22     CLASSES = ["background", "aeroplane", "bicycle", "bird", "boat",
23         "bottle", "bus", "car", "cat", "chair", "cow", "diningtable",
24         "dog", "horse", "motorbike", "person", "pottedplant", "sheep",
25         "sofa", "train", "tvmonitor"]
26     COLORS = np.random.uniform(0, 255, size=(len(CLASSES), 3))
27
28     # load our serialized model from disk
29     print("[INFO] loading model...")
30     net = cv2.dnn.readNetFromCaffe(args["prototxt"], args["model"])

```

We initialize `CLASSES`, our class labels, and corresponding `COLORS`, for on-frame text and bounding boxes (**Lines 22-26**), followed by loading the serialized neural network model (**Line 30**).

Next, we'll initialize the video stream object and frames per second counter:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```

32     # initialize the video stream, allow the camera sensor to warm up,
33     # and initialize the FPS counter
34     print("[INFO] starting video stream...")
35     vs = VideoStream(src=0).start()
36     # vs = VideoStream(usePiCamera=True).start()
37     time.sleep(2.0)
38     fps = FPS().start()

```

We initialize the video stream and allow the camera warm up for 2.0 seconds (**Lines 35-37**).

On **Line 35** we initialize our `VideoStream` using a USB camera. If you are using the Raspberry Pi camera module you'll want to comment out **Line 35** and uncomment **Line 36** (which will enable you to access the Raspberry Pi camera module via the `VideoStream` class).

From there we start our `fps` counter on **Line 38**.

We are now ready to loop over frames from our input video stream:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```

40     # loop over the frames from the video stream
41     while True:
42         # grab the frame from the threaded video stream and resize it

```

```

43         # to have a maximum width of 400 pixels
44         frame = vs.read()
45         frame = imutils.resize(frame, width=400)
46
47         # grab the frame dimensions and convert it to a blob
48         (h, w) = frame.shape[:2]
49         blob = cv2.dnn.blobFromImage(cv2.resize(frame, (300, 300)),
50                                     0.007843, (300, 300), 127.5)
51
52         # pass the blob through the network and obtain the detections and
53         # predictions
54         net.setInput(blob)
55         detections = net.forward()

```

**Lines 41-55** simply grab and resize a `frame`, convert it to a `blob`, and pass the `blob` through the neural network, obtaining the `detections` and bounding box predictions.

From there we need to loop over the `detections` to see what objects were detected in the `frame`:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```

57         # loop over the detections
58         for i in np.arange(0, detections.shape[2]):
59             # extract the confidence (i.e., probability) associated with
60             # the prediction
61             confidence = detections[0, 0, i, 2]
62
63             # filter out weak detections by ensuring the `confidence` is
64             # greater than the minimum confidence
65             if confidence > args["confidence"]:
66                 # extract the index of the class label from the
67                 # `detections`, then compute the (x, y)-coordinates of
68                 # the bounding box for the object
69                 idx = int(detections[0, 0, i, 1])
70                 box = detections[0, 0, i, 3:7] * np.array([w, h, w, h])
71                 (startX, startY, endX, endY) = box.astype("int")
72
73                 # draw the prediction on the frame
74                 label = "{}: {:.2f}%".format(CLASSES[idx],
75                                             confidence * 100)
76                 cv2.rectangle(frame, (startX, startY), (endX, endY),
77                             COLORS[idx], 2)
78                 y = startY - 15 if startY - 15 > 15 else startY + 15
79                 cv2.putText(frame, label, (startX, y),
80                             cv2.FONT_HERSHEY_SIMPLEX, 0.5, COLORS[idx], 2)

```

On **Lines 58-80**, we loop over our `detections`. For each detection we examine the `confidence` and ensure the corresponding probability of the detection is above a predefined threshold. If it is, then we extract the class label and compute  $(x, y)$  bounding box coordinates. These coordinates will enable us to draw a bounding box around the object in the image along with the associated class label.

From there we'll finish out the loop and do some cleanup:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
82         # show the output frame
83         cv2.imshow("Frame", frame)
84         key = cv2.waitKey(1) & 0xFF
85
86         # if the `q` key was pressed, break from the loop
87         if key == ord("q"):
88             break
89
90         # update the FPS counter
91         fps.update()
92
93     # stop the timer and display FPS information
94     fps.stop()
95     print("[INFO] elapsed time: {:.2f}".format(fps.elapsed()))
96     print("[INFO] approx. FPS: {:.2f}".format(fps.fps()))
97
98     # do a bit of cleanup
99     cv2.destroyAllWindows()
100    vs.stop()
```

**Lines 82-91** close out the loop — we show each frame, `break` if 'q' key is pressed, and update our `fps` counter.

The final terminal message output and cleanup is handled on **Lines 94-100**.

Now that our brief explanation of `real_time_object_detection.py` is finished, let's examine the results of this approach to obtain a baseline.

Go ahead and use the ***"Downloads"*** section of this post to download the source code and pre-trained models.

From there, execute the following command:

Raspberry Pi: Deep learning object detection with OpenCV

Shell

```
1 $ python real_time_object_detection.py \
2     --prototxt MobileNetSSD_deploy.prototxt.txt \
3     --model MobileNetSSD_deploy.caffemodel
4 [INFO] loading model...
5 [INFO] starting video stream...
6 [INFO] elapsed time: 54.70
7 [INFO] approx. FPS: 0.90
```

As you can see from my results we are obtaining ~0.9 frames per second throughput using this method and the Raspberry Pi.

Compared to the [6-7 frames per second using our laptop/desktop](#) we can see that the Raspberry Pi is *substantially* slower.

That's not to say that the Raspberry Pi is unusable when applying deep learning object detection, but you need to set your expectations on what's

realistic (even when applying our [OpenCV + Raspberry Pi optimizations](#)).

**Note:** For what it' s worth, I could only obtain **0.49 FPS** when **NOT** using our optimized OpenCV + Raspberry Pi install — that just goes to show you [how much of a difference NEON and VFPV3 can make](#).

## A different approach to object detection on the Raspberry Pi

Using the example from the previous section we see that calling `net.forward()` is a **blocking operation** — the rest of the code in the `while` loop is not allowed to complete until `net.forward()` returns the `detections`.

So, what if `net.forward()` was not a blocking operation?

Would we able to obtain a faster frames per second throughput?

Well, that' s a loaded question.

No matter what, it will take approximately a little over a second for `net.forward()` to complete using the Raspberry Pi and this particular architecture — that cannot change.

But what we *can* do is create a separate process that is solely responsible for applying the deep learning object detector, thereby unblocking the main thread of execution and allow our `while` loop to continue.

Moving the predictions to separate process will give the *illusion* that our Raspberry Pi object detector is running faster than it actually is, when in reality the `net.forward()` computation is still taking a little over one second.

The only problem here is that our output object detection predictions will lag behind what is currently being displayed on our screen. If you detecting *fast-moving objects* you may miss the detection entirely, or at the very least, the object will be out of the frame before you obtain your detections from the neural network.

Therefore, this approach should only be used for *slow-moving objects* where we can tolerate lag.

To see how this multiprocessing method works, open up a new file, name it `pi_object_detection.py`, and insert the following code:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
1 # import the necessary packages
```

```

2 from imutils.video import VideoStream
3 from imutils.video import FPS
4 from multiprocessing import Process
5 from multiprocessing import Queue
6 import numpy as np
7 import argparse
8 import imutils
9 import time
10 import cv2

```

For the code walkthrough in this section, I'll be pointing out and explaining the *differences* (there are quite a few) compared to our non-multiprocessing method.

Our imports on **Lines 2-10** are mostly the same, but notice the imports of `Process` and `Queue` from Python's [multiprocessing package](#).

Next, I'd like to draw your attention to a new function, `classify_frame`:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```

12 def classify_frame(net, inputQueue, outputQueue):
13     # keep looping
14     while True:
15         # check to see if there is a frame in our input queue
16         if not inputQueue.empty():
17             # grab the frame from the input queue, resize it, and
18             # construct a blob from it
19             frame = inputQueue.get()
20             frame = cv2.resize(frame, (300, 300))
21             blob = cv2.dnn.blobFromImage(frame, 0.007843,
22                                         (300, 300), 127.5)
23
24             # set the blob as input to our deep learning object
25             # detector and obtain the detections
26             net.setInput(blob)
27             detections = net.forward()
28
29             # write the detections to the output queue
30             outputQueue.put(detections)

```

Our new `classify_frame` function is responsible for our multiprocessing — later on we'll set it up to run in a child process.

The `classify_frame` function takes three parameters:

- `net`: the neural network object.
- `inputQueue`: our FIFO (first in first out) queue of frames for object detection.
- `outputQueue`: our FIFO queue of detections which will be processed in the main thread.

This child process will loop continuously until the parent exits and effectively terminates the child.

In the loop, if the `inputQueue` contains a `frame`, we grab it, and then pre-process it and create a `blob` (**Lines 16-22**), just as we have done in the previous script.

From there, we send the `blob` through the neural network (**Lines 26-27**) and place the `detections` in an `outputQueue` for processing by the parent.

Now let's parse our command line arguments:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
32 # construct the argument parse and parse the arguments
33 ap = argparse.ArgumentParser()
34 ap.add_argument("-p", "--prototxt", required=True,
35                 help="path to Caffe 'deploy' prototxt file")
36 ap.add_argument("-m", "--model", required=True,
37                 help="path to Caffe pre-trained model")
38 ap.add_argument("-c", "--confidence", type=float, default=0.2,
39                 help="minimum probability to filter weak detections")
40 args = vars(ap.parse_args())
```

There is no difference here — we are simply parsing the same command line arguments on **Lines 33-40**.

Next we initialize some variables just as in our previous script:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
42 # initialize the list of class labels MobileNet SSD was trained to
43 # detect, then generate a set of bounding box colors for each class
44 CLASSES = ["background", "aeroplane", "bicycle", "bird", "boat",
45            "bottle", "bus", "car", "cat", "chair", "cow", "diningtable",
46            "dog", "horse", "motorbike", "person", "pottedplant", "sheep",
47            "sofa", "train", "tvmonitor"]
48 COLORS = np.random.uniform(0, 255, size=(len(CLASSES), 3))
49
50 # load our serialized model from disk
51 print("[INFO] loading model...")
52 net = cv2.dnn.readNetFromCaffe(args["prototxt"], args["model"])
```

This code is the same — we initialize class labels, colors, and load our model.

Here's where things get different:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
54 # initialize the input queue (frames), output queue (detections),
55 # and the list of actual detections returned by the child process
56 inputQueue = Queue(maxsize=1)
57 outputQueue = Queue(maxsize=1)
58 detections = None
```

On **Lines 56-58** we initialize an `inputQueue` of frames, an `outputQueue` of detections, and a `detections` list.



Our `inputQueue` will be populated by the parent and processed by the child — it is the input to the child process. Our `outputQueue` will be populated by the child, and processed by the parent — it is output from the child process. Both of these queues trivially have a size of *one* as our neural network will only be applying object detections to one frame at a time.

Let's initialize and start the child process:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
60 # construct a child process *independent* from our main process of
61 # execution
62 print("[INFO] starting process...")
63 p = Process(target=classify_frame, args=(net, inputQueue,
64     outputQueue,))
65 p.daemon = True
66 p.start()
```

It is very easy to construct a child process with Python's multiprocessing module — simply specify the `target` function and `args` to the function as we have done on **Lines 63 and 64**.

**Line 65** specifies that `p` is a [daemon process](#), and **Line 66** kicks the process off.

From there we'll see some more familiar code:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
68 # initialize the video stream, allow the camera sensor to warmup,
69 # and initialize the FPS counter
70 print("[INFO] starting video stream...")
71 vs = VideoStream(src=0).start()
72 # vs = VideoStream(usePiCamera=True).start()
73 time.sleep(2.0)
74 fps = FPS().start()
```

Don't forget to change your video stream object to use the PiCamera if you desire by switching which line is commented (**Lines 71 and 72**).

Once our `vs` object and `fps` counters are initialized, we can loop over the video frames:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```
76 # loop over the frames from the video stream
77 while True:
78     # grab the frame from the threaded video stream, resize it, and
79     # grab its dimensions
80     frame = vs.read()
81     frame = imutils.resize(frame, width=400)
82     (fH, fW) = frame.shape[:2]
```



```

118         cv2.rectangle(frame, (startX, startY), (endX, endY),
119                        COLORS[idx], 2)
120         y = startY - 15 if startY - 15 > 15 else startY + 15
121         cv2.putText(frame, label, (startX, y),
122                    cv2.FONT_HERSHEY_SIMPLEX, 0.5, COLORS[idx], 2)

```

If our `detections` list is populated (it is not `None`), we loop over the detections as we have done in the previous section's code.

In the loop, we extract and check the `confidence` against the threshold (**Lines 100-105**), extract the class label index (**Line 110**), and draw a box and label on the frame (**Lines 111-122**).

From there in the while loop we'll complete a few remaining steps, followed by printing some statistics to the terminal, and performing cleanup:

Raspberry Pi: Deep learning object detection with OpenCV

Python

```

124         # show the output frame
125         cv2.imshow("Frame", frame)
126         key = cv2.waitKey(1) & 0xFF
127
128         # if the `q` key was pressed, break from the loop
129         if key == ord("q"):
130             break
131
132         # update the FPS counter
133         fps.update()
134
135     # stop the timer and display FPS information
136     fps.stop()
137     print("[INFO] elapsed time: {:.2f}".format(fps.elapsed()))
138     print("[INFO] approx. FPS: {:.2f}".format(fps.fps()))
139
140     # do a bit of cleanup
141     cv2.destroyAllWindows()
142     vs.stop()

```

In the remainder of the loop, we display the frame to the screen (**Line 125**) and capture a key press and check if it is the quit key at which point we break out of the loop (**Lines 126-130**). We also update our `fps` counter.

To finish out, we stop the `fps` counter, print our time/FPS statistics, and finally close windows and stop the video stream (**Lines 136-142**).

Now that we're done walking through our new multiprocessing code, let's compare the method to the single thread approach from the previous section.

Be sure to use the ***"Downloads"*** section of this blog post to download the source code + pre-trained MobileNet SSD neural network. From there, execute the following command:

Shell

```
1 $ python pi_object_detection.py \  
2     --prototxt MobileNetSSD_deploy.prototxt.txt \  
3     --model MobileNetSSD_deploy.caffemodel  
4 [INFO] loading model...  
5 [INFO] starting process...  
6 [INFO] starting video stream...  
7 [INFO] elapsed time: 48.55  
8 [INFO] approx. FPS: 27.83
```

Here you can see that our `while` loop is capable of processing 27 frames per second. However, this throughput rate is an illusion — the neural network running in the background is still only capable of processing 0.9 frames per second.

**Note:** *I also tested this code on the Raspberry Pi camera module and was able to obtain 60.92 frames per second over 35 elapsed seconds.*

The difference here is that we can obtain real-time *throughput* by displaying each new input frame in real-time and then drawing any previous `detections` on the current frame.

Once we have a new set of `detections` we then draw the new ones on the frame.

This process repeats until we exit the script. **The downside is that we see substantial lag.** There are clips in the above video where we can see that all objects have clearly left the field of view...

...however, our script still reports the objects as being present.

Therefore, you should ***consider only using this approach when:***

1. Objects are slow moving and the previous detections can be used as an approximation to the new location.
2. Displaying the actual frames themselves in real-time is paramount to user experience.

## Summary

In today's blog post we examined using the Raspberry Pi for object detection using deep learning, OpenCV, and Python.

As our results demonstrated we were able to get up to **0.9 frames per second**, which is not fast enough to constitute real-time detection. That said,

given the limited processing power of the Pi, 0.9 frames per second is still reasonable for some applications.

We then wrapped up this blog post by examining an alternate method to deep learning object detection on the Raspberry Pi by using multiprocessing. Whether or not this second approach is suitable for you is again highly dependent on your application.

If your use case involves low traffic object detection where the objects are slow moving through the frame, then you can certainly consider using the Raspberry Pi for deep learning object detection. However, if you are developing an application that involves many objects that are fast moving, you should instead consider faster hardware.

Thanks for reading and enjoy!

And if you're interested in studying deep learning in more depth, be sure to take a look at my new book, [\*Deep Learning for Computer Vision with Python\*](#).

Whether this is the first time you've worked with machine learning and neural networks or you're already a seasoned deep learning practitioner, my new book is engineered from the ground up to help you reach expert status.

[\*\*Just click here to start your journey to deep learning mastery.\*\*](#)