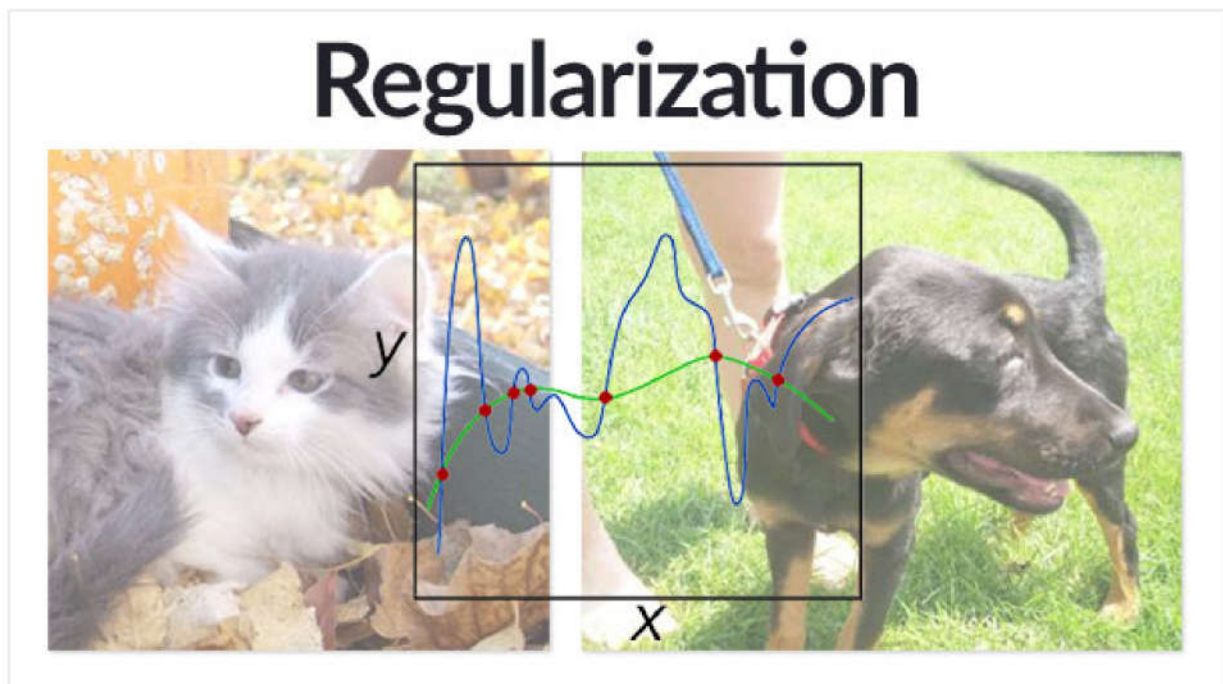




# Understanding regularization for image classification and machine learning

by **Adrian Rosebrock** on September 19, 2016 in **Deep Learning, Machine Learning, Tutorials**



In previous tutorials, I've discussed two important loss functions: *Multi-class SVM loss* and *cross-entropy loss* (which we usually refer to in conjunction with Softmax classifiers).

In order to keep our discussions of these loss functions straightforward, I purposely left out an important component: **regularization**.

While our loss function allows us to determine how well (or poorly) our set of parameters (i.e., weight matrix, and bias vector) are performing on a given classification task, the loss function itself does not take into account how the weight matrix “looks”.

What do I mean by “looks”?

Well, keep in mind that there may be an *infinite* set of parameters that obtain reasonable classification accuracy on our dataset — how do we go about choosing a set of

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

Click to answer

There are various types of regularization techniques, such as L1 regularization, L2 regularization, and Elastic Net — and in the context of Deep Learning, we also have *dropout* (although dropout is more-so a *technique* rather than an actual *function*).

Inside today's tutorial, we'll mainly be focusing on the former rather than the later. Once we get to more advanced deep learning tutorials, I'll dedicate time to discussing dropout as well.

In the remainder of this blog post, I'll be discussing regularization further. I'll also demonstrate how to update our Multi-class SVM loss and cross-entropy loss functions to include regularization. Finally, we'll write some Python code to construct a classifier that applies regularization to an image classification problem.

**Looking for the source code to this post?**  
**[Jump right to the downloads section.](#)**

## Understanding regularization for image classification and machine learning

The remainder of this blog post is broken into four parts. First, we discuss what regularization is. I then detail how to update our loss function to include the regularization term.

From there, I list out three common types of regularization you'll likely see when performing image classification and machine learning, *especially* in the context of neural networks and deep learning.

Finally, I'll provide a Python + scikit-learn example that demonstrates how to apply regularization to an image classification dataset.

### What is regularization and why do we need it?

**Regularization helps us tune and control our model complexity**, ensuring that our models are better at making (correct) classifications — or more simply, *the ability to generalize*.

If we don't apply regularization, our classifiers can easily become too complex and *overfit* to our training data, in which case we lose the ability to generalize to our testing data (and data points outside the testing set as well).

Similarly, without applying regularization we also run the risk of *underfitting*. In this case, our model performs poorly on the training our — our classifier is not able to model the relationship between the input data and the output class labels.

Underfitting is relatively easy to catch — you examine the classification accuracy on your training data and take a look at your model.

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

[Click to answer](#)

## Overfitting is a different beast entirely though.

While you can certainly monitor your training accuracy and recognizing when your classifier is performing *too well* on the training data and *not good enough* on the testing data, it becomes harder to correct.

There is also the problem that you can walk a very fine line between model complexity — if you simplify your model *too much*, then you'll be back to underfitting.

A better approach is to apply *regularization* which will help our model generalize and lead to less overfitting.

The best way to understand regularization is to see the implications it has on our loss function, which I discuss in the next section.

## Updating our loss function to include regularization

Let's start with our [Multi-class SVM loss function](#):

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

The loss for the entire training set can be written as:

$$L = \frac{1}{N} \sum_{i=1}^N L_i$$

Now, let's say that we have obtained a weight matrix  $W$  such that *every data point* in our training set is classified 100% correctly — this implies that our loss  $L = 0$  for all  $L_i$ .

Awesome, we're getting 100% accuracy — but let me ask you a question about this weight matrix  $W$  — **is this matrix unique?**

**Or in other words, are there BETTER choices of  $W$  that will improve our model's ability to generalize and reduce overfitting?**

If there is such a  $W$ , how do we know? And how can we incorporate this type of penalty into our loss function?

The answer is to define a **regularization penalty**, a function that operates on our weight matrix  $W$ .

The regularization penalty is commonly written as the function  $R(W)$ .

Below is the most common regularization penalty, L2 regularization:

$$R(W) = \sum_i \sum_j W_{i,j}^2$$

What is this function doing exactly?

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

[Click to answer](#)



```

1 penalty = 0
2
3 for i in np.arange(0, W.shape[0]):
4     for j in np.arange(0, W.shape[1]):
5         penalty += (W[i][j] ** 2)

```

What we are doing here is looping over all entries in the matrix and taking the sum of squares. There are more efficient ways to compute this of course, I'm just simplifying the code as a matter of explanation.

The sum of squares in the L2 regularization penalty discourages large weights in our weight matrix  $W$ , preferring smaller ones.

Why might we want to discourage large weight values?

In short, by penalizing large weights we can improve our ability to generalize, and thereby reduce overfitting.

Think of it this way — the larger a weight value is, the more influence it has on the output prediction. This implies that dimensions with larger weight values can almost singlehandedly control the output prediction of the classifier (provided the weight value is large enough, of course), which will almost certainly lead to overfitting.

To mitigate the affect various dimensions have on our output classifications, we apply regularization, thereby seeking  $W$  values that take into account *all* of the dimensions rather than the few with large values.

In practice, you may find that regularization hurts your training accuracy slightly, but actually *increases your testing accuracy* (your ability to generalize).

Again, our loss function has the following basic form, but now we just add in regularization:

$$L = \frac{1}{N} \sum_{i=1}^N L_i + \lambda R(W)$$

The first term we have already seen before — this is the average loss over all samples in our training set.

**The second term is new — *this is our regularization term.***

The  $\lambda$  variable is a hyperparameter that controls the *amount* or *strength* of the regularization we are applying. In practice, both the learning rate  $\alpha$  and regularization term  $\lambda$  are hyperparameters that you'll spend most of your time tuning.

Expanding the Multi-class SVM loss to include regularization yields the final equation:

$$L = \frac{1}{N} \sum_{i=1}^N \sum_{j \neq y_i} [\max(0, s_j - s_{y_i} + 1)] + \lambda \sum_i \sum_j W_{i,j}^2$$

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

Click to answer

For a more mathematically motivated discussion of regularization, take a look at [Karpathy's excellent slides](#) from the CS231n course.

## Types of regularization techniques

In general, you'll see three common types of regularization.

The first, we reviewed earlier in this blog post, L2 regularization;

$$R(W) = \sum_i \sum_j W_{i,j}^2$$

We also have L1 regularization which takes the absolute value rather than the square:

$$R(W) = \sum_i \sum_j |W_{i,j}|$$

Elastic Net regularization seeks to combine both L1 and L2 regularization:

$$R(W) = \sum_i \sum_j \beta W_{i,j}^2 + |W_{i,j}|$$

In terms of which regularization method you should be using (including none at all), you should treat this choice as a hyperparameter you need to optimize over and perform experiments to determine *if* regularization should be applied, and if so, *which method* of regularization.

Finally, I'll note that there is another *very common* type of regularization that we'll see in a future tutorial — **dropout**.

Dropout is frequently used in Deep Learning, especially with Convolutional Neural Networks.

Unlike L1, L2, and Elastic Net regularization, which boil down to functions defined in the form  $R(W)$ , dropout is an actual *technique* we apply to the connections between nodes in a Neural Network.

As the name implies, connections “dropout” and randomly disconnect during training time, ensuring that no one node in the network becomes fully responsible for “learning” to classify a particular label. I'll save a more thorough discussion of dropout for a future blog post.

## Image classification using regularization with Python and scikit-learn

Now that we've discussed regularization in the context of machine learning, let's look at some code that actually *performs* various types of regularization.

All of the code associated with this blog post, except for the final code block, has already been reviewed extensively in previous blog posts in this series.

Therefore, for a thorough review of the actual process used to extract features and construct the training and testing split for the Kaggle Dogs vs. Cats dataset, I'll refer you to the [introduction to linear classification tutorial](#).

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

[Click to answer](#)

The code block below demonstrates how to apply the Stochastic Gradient Descent (SGD) classifier with log-loss (i.e., Softmax) and various types of regularization methods to our dataset:

Understanding regularization for image classification and machine le	Python
<pre> 73 # loop over our set of regularizers 74 for r in (None, "l1", "l2", "elasticnet"): 75     # train a Stochastic Gradient Descent classifier using a softmax 76     # loss function, the specified regularizer, and 10 epochs 77     print("[INFO] training model with `{}` penalty".format(r)) 78     model = SGDClassifier(loss="log", penalty=r, random_state=967, 79         n_iter=10) 80     model.fit(trainData, trainLabels) 81 82     # evaluate the classifier 83     acc = model.score(testData, testLabels) 84     print("[INFO] `{}` penalty accuracy: {:.2f}%".format(r, acc * 100)) </pre>	

On **Line 74** we start looping over our regularization methods, including `None` for *no regularization*.

We then train our `SGDClassifier` on **Lines 78-80** using the specified regularization method.

**Lines 83 and 84** evaluate our trained classifier on the testing data and display the accuracy.

Below I have included a screenshot from executing the script on my machine:

```

[INFO] constructing training/testing split...
[INFO] training model with `None` penalty
[INFO] `None` penalty accuracy: 63.58%
[INFO] training model with `l1` penalty
[INFO] `l1` penalty accuracy: 64.02%
[INFO] training model with `l2` penalty
[INFO] `l2` penalty accuracy: 64.38%
[INFO] training model with `elasticnet` penalty
[INFO] `elasticnet` penalty accuracy: 64.40%
(keras) adrian@annalee:~/pyimagesearch/pyimg-at-posts/regularization$

```

Figure 1: Applying no regularization, L1 regularization, L2 regularization, and Elastic Net regularization to our classification project.

As we can see, classification accuracy on the testing set *improves* as regularization is introduced.

We obtain **63.58%** accuracy with no regularization. Applying L1 regularization increases our accuracy to **64.02%**. L2 regularization improves again to **64.38%**. Finally, Elastic Net, which combines both L1 and L2 regularization obtains the highest accuracy of **64.40%**.

Does this mean that we should *always* apply Elastic Net regularization?

Of course not — this is entirely dependent on your dataset and features. You should treat regularization, and any parameters associated with your regularization method, as **hyperparameters that need to be searched over**.

## Summary

Adrian Rosebrock, Chief PyImageSearcher

What are you most interested in studying right now?

Click to answer



Regularization works by examining our weight matrix  $W$  and penalizing it if it does not confirm to the specified penalty function.

Applying this penalty helps ensure we learn a weight matrix  $W$  that generalizes better and thereby helps lesson the negative affects of overfitting.

In practice, you should apply hyperparameter tuning to determine:

1. If regularization should be applied, and if so, which regularization method should be used.
2. The strength of the regularization (i.e., the  $\lambda$  variable).

You may notice that applying regularization may actually *decrease* your training set classification accuracy — this is acceptable provided that your testing set accuracy *increases*, which would be a demonstration of regularization in action (i.e., avoiding/lessening the impact of overfitting).

In next week's blog post, I'll be discussing how to build a simple feedforward neural network using Python and Keras. ***Be sure to enter your email address in the form below to be notified when this blog post goes live!***

## Downloads:



If you would like to download the code and images used in this post, please enter your email address in the form below. Not only will you get a .zip of the code, I'll also send you a **FREE 17-page Resource Guide on Computer Vision, OpenCV, and Deep Learning**. Inside you'll find my hand-picked tutorials, books, courses, and libraries to help you master CV and DL! Sound good? If so, enter your email address and I'll send you the code immediately!

Email address:

DOWNLOAD THE CODE!

### Resource Guide (it's totally free).


Enter your email address below to get my **free 17-page Computer Vision, OpenCV, and Deep Learning Resource Guide PDF**. Inside you'll find my hand-picked tutorials, books, courses, and Python libraries to help you master computer vision and deep learning!

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

Click to answer



🔖 **classification, deep learning, dogs and cats, elastic net, l1-norm, l2-norm, machine learning, regularization**

< [Softmax Classifiers Explained](#)

[A simple neural network with Python and Keras](#) >

## 6 Responses to *Understanding regularization for image classification and machine learning*



**Ashis Samal** September 20, 2016 at 1:42 pm #

REPLY ↩

Thank you very much. Nice explanation .



**Adrian Rosebrock** September 21, 2016 at 2:11 pm #

REPLY ↩

Thanks Ashis, I'm glad you found it helpful!



**Manuel** May 25, 2017 at 3:40 pm #

REPLY ↩

Hi Adrian, a question regarding the search for the best regularization term using the techniques you explain in your "How to tune hyperparameters with Python and scikit-learn" post.

How should I specify in the RandomizedSearchCV function, that the score I want to maximize is the

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

[Click to answer](#)



**Manuel** May 25, 2017 at 4:06 pm #

REPLY ↩

Wait, I think that CV at the end of the function means cross-validation, so accuracy report is already being done on a cross-validated set, which would be equivalent to report on a validation set, right?

**Adrian Rosebrock** May 28, 2017 at 1:22 am #

REPLY ↩

The cross-validation will be performed by dividing your training set into N parts, training on all but one of the sets, and then validating on the other. This is done for each data split and serves as a proxy for your validation.

**Igor** June 15, 2018 at 12:55 pm #

REPLY ↩

Hi Adrian,

I understood that regularization helps to fix the overfitting. But, you said "... without applying regularization we also run the risk of underfitting." Why is that? Regularization makes the network simpler to avoid overfitting, so how does regularization help with underfitting?

## Before you leave a comment...

Hey, Adrian here, author of the PyImageSearch blog. I'd love to hear from you, but before you submit a comment, **please follow these guidelines:**

- **If you have a question, read the comments first.** You should also search this page (i.e., *ctrl + f*) for keywords related to your question. It's likely that I have already addressed your question in the comments.
- **If you are copying and pasting code/terminal output, please don't.** Reviewing another programmers' code is a very time consuming and tedious task, and due to the volume of emails and contact requests I receive, I simply cannot do it.
- **Be respectful of the space.** I put a *lot* of my own personal time into creating these free weekly tutorials. On average, each tutorial takes me 15-20 hours to put together. I love offering these guides to you and I take pride in the content I create. Therefore, I will not approve comments that include large code blocks/terminal output as it destroys the formatting of the page. Kindly be respectful of this space.
- **Be patient.** I receive 200+ comments and emails per day. Due to spam, and my desire to personally answer as many questions as I can, I hand moderate all new comments (typically once per week). I try to answer as many questions as I can, but I'm only one person. Please don't be

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

Click to answer

left over, I focus on the community at large and attempt to answer as many of those questions as I possibly can.

Thank you for keeping these guidelines in mind before submitting your comment.

## Leave a Reply

Name (required)

Email (will not be published) (required)


Website

SUBMIT COMMENT

Search...



## Resource Guide (it's totally free).



Get your **FREE 17 page Computer Vision, OpenCV, and Deep Learning Resource Guide PDF**. Inside you'll find my hand-picked tutorials, books, courses, and libraries to help you master CV and DL.

[Download for Free!](#)

## Raspberry Pi for Computer Vision

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

[Click to answer](#)

# KICKSTARTER



You can teach your **Raspberry Pi** to “see” using **Computer Vision**, **Deep Learning**, and **OpenCV**. **Let me show you how.**

[CLICK HERE TO LEARN MORE](#)

**Deep Learning for Computer Vision with Python Book — OUT NOW!**




You're interested in deep learning and computer vision, *but you don't know how to get started*. Let me help. **My new book will teach you all you need to know about deep learning.**

Adrian Rosebrock, Chief PyImageSearcher

**What are you most interested in studying right now?**

[Click to answer](#)





Are you interested in **detecting faces in images & video**? But **tired of Googling for tutorials** that *never work*? Then let me help! I guarantee that my new book will turn you into a **face detection ninja** by the end of this weekend. [Click here to give it a shot yourself.](#)

[CLICK HERE TO MASTER FACE DETECTION](#)

**PyImageSearch Gurus: NOW ENROLLING!**

The **PyImageSearch Gurus** course is **now enrolling!** Inside the course you'll learn how to perform:

- Automatic License Plate Recognition (ANPR)
- Deep Learning
- Face Recognition
- *and much more!*

**Click the button below to learn more about the course, take a tour, and get 10 (FREE) sample lessons.**

[TAKE A TOUR & GET 10 \(FREE\) LESSONS](#)

Hello! I'm Adrian Rosebrock


Adrian Rosebrock, Chief PyImageSearcher

**What are you most interested in studying right now?**

[Click to answer](#)



Learn computer vision in a single weekend.



Want to learn computer vision & OpenCV? I can teach you in a **single weekend**. I know. It sounds crazy, but it's no joke. My new book is your **guaranteed, quick-start guide** to becoming an OpenCV Ninja. So why not give it a try? [Click here to become a computer vision ninja.](#)

[CLICK HERE TO BECOME AN OPENCV NINJA](#)

## Subscribe via RSS



**Never miss a post!** Subscribe to the PyImageSearch RSS Feed and keep up to date with my image search engine tutorials, tips, and tricks

### POPULAR

**Raspbian Stretch: Install OpenCV 3 + Python on your Raspberry Pi**  
SEPTEMBER 4, 2017

**Install guide: Raspberry Pi 3 + Raspbian Jessie + OpenCV 3**  
APRIL 18, 2016

**Face recognition with OpenCV, Python, and deep learning**  
JUNE 18, 2018

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

[Click to answer](#)

**Real-time object detection with deep learning and OpenCV**

SEPTEMBER 18, 2017

**Ubuntu 16.04: How to install OpenCV**

OCTOBER 24, 2016

Find me on [Twitter](#), [Facebook](#), and [LinkedIn](#).

[Privacy Policy](#)

© 2019 PyImageSearch. All Rights Reserved.

Adrian Rosebrock, Chief PyImageSearcher



**What are you most interested in studying right now?**

[Click to answer](#)